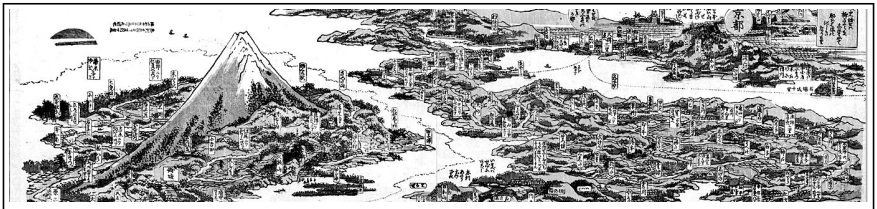
MORNING
UPDATE
TORONTO

The Long View

10° · 2°
CLOUDY, THEN
RAIN

Printed once a day. Finished by breakfast.

WEDNESDAY, NOVEMBER 11, 2026 NO. 2 20 PAGES · 5 SHEETS · $1.00



The old road from Edo to Kyoto, drawn by Katsushika Hokusai around 1818. Today's back page is the second stop on that road, Shinagawa.

Your GIS Assignment 3 is due at 10:59 this morning, and the Practical 7 quiz closes at 7 tonight. An AI Safety Week session at 11:15 overlaps your GGRA30 lecture, so pick one. Expect a cloudy day at about 10°, and real rain on Thursday. Two PHLB18 items are still late.

TODAY · WEDNESDAY

10:59	GGRA30 Assignment 3 due (PostgreSQL in QGIS) Quercus upload · 5%
11:00	GGRA30 lecture HL B101 · until 1:00 · clashes with next
11:15	*AI Safety Week, Day 3: Power Concentration Risk* Design Exchange, 234 Bay St · until 12:15
1:00	ENGA10 lecture (Mrs Dalloway) AC 223 · until 3:00
4:00	GGRA30 practical BV 469 · until 5:00
6:00	*Hinton Lectures, night 3* Livestream · until 7:30

TODAY	THU	FRI	SAT	SUN
10° / 2°	7° / 0°	6° / 2°	6° / 2°	8° / 5°
45%	82% rain			

INSIDE, IN ONE MINUTE

ESSAYS How to take criticism; why India is doing well.

WORLD Indonesia's vice president; Spain's strike; Ebola in Congo.

TO DO

ACADEMICS

- [x] GGRA30 Assignment 3 is due 10:59 this morning. Practical 7 quiz closes 7 tonight. · Todoist
- [ ] GGRA30 midterm grades were posted in Quercus last night. · Quercus mail
- [x] PHLB18 reading response 6 is late. It earns half marks until Nov 16, then nothing. · Todoist
- [ ] PHLB18 Scaffolding 2 is late. It loses 20% for every 3 days. It stays open to Dec 8. · Todoist
- [ ] MGEA01 Achieve quiz, Chapter 11: Thursday 7 pm. · Todoist

TAISI

- [ ] Upload your receipts to HCB, if you haven't. · #taisi-team
- [ ] Shivam's cohort, Mon Nov 16: order the kosher meal too. · Todoist

MISC

- [x] Pay off the credit card by Thursday. · Todoist
- [ ] Read the UTSC Health and Wellness message waiting in Medeo. · UofT mail

FOR YOU

Your Desk

What needs you, and what TAISI talked about

NEEDS YOU

1. **School deadlines today.** GGRA30 Assignment 3 is due at 10:59 this morning, before your lecture ends. The Practical 7 quiz opens at 1:00 and closes at 7:00 tonight. Both are in Todoist. — Todoist · QUERCUS_CONTEXT
2. **Two PHLB18 items are late.** Reading response 6 gets half marks until Nov 16, then is not accepted. Scaffolding 2 loses 20% for every 3 days it is late, so handing it in soon costs less. It closes Dec 8. — Todoist · QUERCUS_CONTEXT
3. **Two course windows close this weekend.** BlueDot's governance course closes Sunday, Nov 15. CLR's course closes Monday, Nov 16. Both are below. — bluedot.org · longtermrisk.org

THIS WEEK AND BEYOND

Wed	AI Safety Week, Day 3: Forecasting AI-Enabled Power Concentration Risk 11:15–12:15 · Design Exchange · overlaps GGRA30 lecture
Wed	Hinton Lectures, night 3, then the Q&A 6:00–7:30 pm · livestream
Thu	MGEA01 lecture, Perfect Competition; Achieve quiz Ch. 11 due 7 pm 7:00–10:00 pm · IC 130
Thu	Credit card payment your Todoist date
Fri	Apart epistemics sprint starts; AI Safety Camp project list goes up online · see Opportunities
Sat	*Marches for an AI treaty* cities in the US, Canada and France · noted in Astral Codex Ten
Sun	COPB50 Informational Interview due; BlueDot closes Nov 15
Nov 16	PHLB18 lecture 10–12 (IA 2150) and tutorial 1–2 (HW 402); reading response 7 due; CLR closes; Shivam's cohort food Monday

YOUR EVENTS (FROM LUMA)

Nov 11 Wed	**AI Safety Week, Day 3.** Power Concentration Risk, 11:15 to 12:15. APPROVED
Nov 9–11 6 pm	**The Hinton Lectures 2026.** Paul Christiano, three evenings. Monday's and Tuesday's recordings are online. REGISTERED

TAISI SLACK · SINCE MONDAY

#taisi-team · money and food

- Isabel reminded everyone to upload receipts to HCB. A locked card hurts people who use the Uber Business account. Joseph said he is working on it.
- You ordered shawarma from Ali Baba for Shivam's cohort, and kosher food from Dr. Laffa for Daniel. Isabel thanked you.

#taisi-team · a study-group offer

- Someone outside TAISI wants to run a linear algebra study group. Joseph thinks it is slightly good, but he would not run it and would not call it a TAISI program. He is happy to refer people and may call the organizer first.
- Boyan said it should not overlap with his reading group. He asked who the group is for, since it says it is for people outside a university program.

Other channels

- No new messages in #writing-desk, #si-main or #hackathons. #general had only a note about spare Hinton tickets for Tuesday.

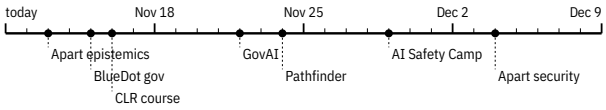
DON'T MISS

Opportunities

Sorted by deadline

Picked for a second-year CS and Public Policy student on TAISI's operations team. ★ = strong fit. Each item says why. Dates were last checked on each program's own page on Nov 9. BlueDot and CLR were checked again today.

THE NEXT FOUR WEEKS



CLOSING SOON

Nov 13
starts Fri

Apart: AI × Epistemics sprint Apart Research

A free online weekend sprint, Nov 13 to 15. Teams build a test, dataset or tool about how AI changes what people believe. $2,000 in prizes. Entries are due Sunday night.

This weekend. No AI safety background is needed, and you can work at your own pace. A low-cost way to get a first research write-up.

HACKATHON ONLINE FREE apartresearch.com/sprints

Nov 15 Sun; closes in 4 days

★ **Frontier AI Governance course (online)** BlueDot Impact

BlueDot's own free online governance course, about 40 hours. Take it as a six-day intensive or over six weeks. Groups run in many time zones.

The governance half of what you did at TAISI, with no travel. No earlier BlueDot course is required. You give your free hours after you apply.

COURSE ONLINE GOVERNANCE FREE bluedot.org/courses/ai-governance

---

Nov 16 Mon, 23:59 GMT

★ **CLR Fundamentals Program** Center on Long-Term Risk

An online reading course, Nov 20 to Dec 21, on how advanced AI could cause suffering on a huge scale. About five hours a week. Meeting times fit the group.

Online, short, and new topics beyond what TAISI's fellowship covered. It runs into exam season, so five hours a week needs a plan.

COURSE ONLINE BEGINNER OK longtermrisk.org/intro-seminar

Nov 22 Sun, 23:59 UK time

★ **GovAI Summer Fellowship 2027 (London)** Centre for the Governance of AI

Three months of AI policy work in London, June 7 to Aug 27, 2027, paid £14,000. There is a research track and an applied track. No degree is required.

A stretch for a second-year, but CS plus policy is exactly their field. GovAI can sponsor UK work visas. Skip the DC version: it needs the right to work in the US.

FELLOWSHIP PAID GOVERNANCE STRETCH governance.ai/opportunities

---

Nov 24 Tue, 23:59 anywhere on Earth

★ **Kairos Pathfinder Fellowship, Spring 2027** Kairos

Mentoring, workshops and group funding for students who organize AI safety or AI policy groups at a university. Online. Decisions come Dec 2 and it starts Dec 5.

You help run TAISI, so you can apply. Each organizer applies alone. Kairos now funds only groups with an organizer in this program, so tell TAISI's leads. Undergrads get no personal stipend.

FIELD-BUILDING OPS REMOTE pathfinder.kairos-project.org

Nov 29 Sun; opens Nov 13

★ **AI Safety Camp 12** AI Safety Camp

An online research program, Jan 9 to Apr 18, 2027. You join a small team on a project picked in advance. The project list goes up Nov 13, when team applications open.

Part-time, remote and open to beginners. A good first research project, and it does not depend on SPAR. Read the list on Friday and apply to the governance or security projects.

RESEARCH REMOTE BEGINNER OK aisafety.camp

Dec 4 starts Fri

★ **Apart: Threat Intelligence & Threat Modelling hackathon** Apart Research

A free online weekend hackathon, Dec 4 to 6, run with Security Level 5. Teams test whether AI can speed up security planning for AI data centres. $5,000 in prizes.

The closest match to your security interest on this page. Entries are due Dec 6, 11:59pm anywhere on Earth. It falls near exams, so decide early.

HACKATHON SECURITY ONLINE FREE apartresearch.com

OPENING SOON: BEST FITS

Dec opens; date not announced

★ **SPAR, Spring 2027** SPAR (run by Kairos)
Remote, part-time research with a mentor, usually 5 to 20 hours a week. Open to undergrads. It has policy projects too. Applications open around December.
Your best next step: real research that fits around classes. Press "Express interest" on the site now so you get the email when it opens.
RESEARCH · REMOTE · UNDERGRAD OK · sparai.org

Dec 11 starts Fri

Apart: Agent Attribution Hackathon Apart Research
A free online weekend sprint, Dec 11 to 13, on telling rogue AI agents apart from ones a human is steering. Entries are due Dec 13.
An AI security topic close to your TryHackMe work. Some other sites still show Nov 27 to 29. Apart's own page says Dec 11.
HACKATHON · SECURITY · ONLINE · FREE · apartresearch.com

TBA not open yet

★ **Vector Institute research internship, summer 2027** Vector Institute
A paid research internship in Toronto, May to August, with a Vector professor. Vector takes summer applications from November to January. The page says they are closed for now.
In Toronto and open to undergrads at any school, in Canada or abroad. Letters are optional for undergrads. You need some machine learning experience.
RESEARCH · TORONTO · PAID · vectorinstitute.ai

TBA likely Feb; not announced

★ **UTSC summer research award (UTEA)** University of Toronto Scarborough
A $7,500 award to do 14 to 16 weeks of summer research with a UTSC professor. The 2027 date is not posted. The 2026 deadline was Feb 27.
On your own campus. Foreign students with a valid study permit for the summer can hold a UTEA. You need a professor to back you, so start asking in November.
RESEARCH · TORONTO · UNDERGRAD OK · PAID · utsc.utoronto.ca/research/utea-ursa

FOR TAISI, ANY TIME

any time

Money for TAISI projects always open
BlueDot Rapid Grants give up to $20,000 for AI safety projects. Manifund keeps a list of who else is funding. It says the Long-Term Future Fund has stopped taking new applications.
Useful for ops: hackathon food, room costs, compute keys. For TAISI's regular budget, the route is now the Pathfinder Fellowship above, due Nov 24.
GRANTS · FOR GROUPS · bluedot.org

A NOTE ON ACCURACY

Programs that don't fit you yet are listed in this morning's email. Always check dates on the program's own site.

NEWS

The World

The facts first. The opinions get labels.



REGION ONE

Indonesia

A House move against the vice president, and a China trade plan.

POLITICS

Every party in Indonesia's House backs a move to impeach the vice president

All party groups in the House of Representatives have handed in a formal motion to start impeachment of Vice President Gibran Rakabuming Raka. It is about his high school papers. The motion is called the Right to Express an Opinion. It is the first step toward asking whether he can stay in office.

Parties that signed include PDI-P, Golkar, NasDem, PKS, the Democratic Party, PKB and Prabowo's own party, Gerindra. A lawmaker from PAN said earlier court rulings show Gibran's school certificate has no legal force. The House leaders say the motion goes to a full meeting of the House on Thursday.

> **Not everyone agrees.** Gibran has said he wants a chance to prove the certificate is real through a court process. Parties in the House say the question is about whether he meets the rules to hold the job.

WHY IT MATTERS: *On Tuesday the palace said Gibran wanted to resign. Now the House is moving to decide the matter itself.*

ANTARA · Nov 11

TRADE

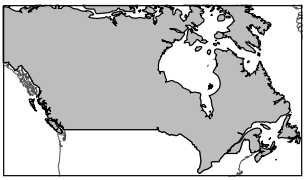
Prabowo and Xi plan to talk about a China–Indonesia trade deal

President Prabowo Subianto and China's president, Xi Jinping, will meet at the APEC summit in Shenzhen. APEC is a group of countries around the Pacific Ocean that talk about trade. A foreign ministry official said one point will be an agreement to start talks on a trade deal called a CEPA, short for Comprehensive Economic Partnership Agreement.

Prabowo speaks to the APEC CEO Summit on Nov 18. The leaders' meeting is on Nov 19. The official said the schedule is packed, so Indonesia may hold only short, informal chats with other partners.

WHY IT MATTERS: *A trade deal with China, Indonesia's biggest trading partner, would shape prices and jobs.*

ANTARA · Nov 11



REGION TWO

Canada

A meeting with China's leader, and a rate rise that may come.

CHINA

Carney will raise foreign interference when he meets Xi

Prime Minister Mark Carney said on Tuesday that he will bring up foreign interference when he meets China's president, Xi Jinping, next week. The two are expected to meet at the APEC summit in Shenzhen.

Carney also said the two countries' partnership is growing but has "clear guardrails". A public inquiry found in January 2025 that China was the top source of foreign interference in Canadian democracy. Carney said a new ambassador to China will be named soon. The post has been empty since spring.

WHY IT MATTERS: *Canada is rebuilding trade ties with China while worrying about its meddling.*

Global News (Canadian Press) · Nov 10

RATES

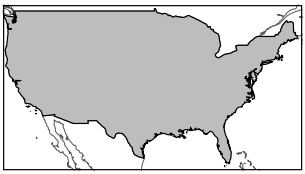
The Bank of Canada says it may raise interest rates again, but slowly

The bank's summary of its last meeting says members will decide "one meeting at a time". Last month it raised its key rate by a quarter point to 2.5%. That was its first rise since 2023.

It fears high oil prices, because of the war in the Middle East, could spread to other prices. It expects inflation to reach 3.5% by the end of this year. The next decision is Dec 9.

WHY IT MATTERS: *Interest rates set what Canadians pay on mortgages and loans.*

Global News (Canadian Press) · Nov 10



REGION THREE

United States

A House vote on the defense secretary, prices still rising, and the Justice Department dropping cases.

CONGRESS

The House sends the effort to impeach Defense Secretary Hegseth back to committee

On Tuesday the House voted to send the impeachment articles back to a committee. That effectively ends the effort for now. Representative Thomas Massie of Kentucky wrote them. He lost his primary. The articles accuse Pete Hegseth of illegal wars and unlawful strikes.

Hegseth said the charges are policy disagreements, and he would carry out the same operations again. President Trump praised him. The same day, House Republicans chose Jim Jordan as their leader, and Senate Republicans kept John Thune.

> **Not everyone agrees.** Massie says Hegseth is a "threat to the Constitution". Hegseth calls it a "political spectacle". Don Bacon, a Republican, has doubts about Hegseth but opposed the impeachment as written.

WHY IT MATTERS: *It shows how much Republicans will back the administration after losing the House.*

PBS NewsHour · Nov 10

PRICES

US prices rose 0.4% in October, and about 4% a year is expected

The latest Consumer Price Index showed prices up four-tenths of a percent in a month, more than expected. Food and energy led the rise. Housing rose 0.3%.

Economists expect about 4% inflation this year, faster than pay is rising. Austan Goolsbee, who runs the Federal Reserve Bank of Chicago, called the report "distressing". He said price rises in services, such as health care, are not mainly from oil or tariffs.

WHY IT MATTERS: *Prices are rising faster than pay, so most households can buy less.*

PBS NewsHour · Nov 10

JUSTICE

The Justice Department asks to wipe out Steve Bannon's contempt conviction

Steve Bannon, an adviser to President Trump, was found guilty of contempt of Congress after he ignored a subpoena from the House committee that studied the Jan. 6 attack on the Capitol. A jury took less than three hours, and an appeals court kept the verdict.

Now the Justice Department has asked a court to drop the case "in the

interests of justice." Bannon's lawyers do not oppose it. Legal experts told NPR it follows a pattern. In April the department asked to cancel the convictions of about a dozen leaders of the Oath Keepers and Proud Boys groups.

> **Not everyone agrees.** Bruce Searby, a defense lawyer and former prosecutor, says the courts hear no opposing side. The Justice Department did not answer NPR's request for comment.

WHY IT MATTERS: *Courts normally hear two sides. Here, no one is arguing against the request.*

NPR · Nov 11

REGION FOUR

Around the World

Peace talks in Abu Dhabi, a general strike in Spain, Ebola in Congo, and floods feared in Kenya.

UKRAINE

Ukraine, Russia and the US hold a first day of peace talks in Abu Dhabi

Senior officials met for about four hours on Tuesday and will meet again today. They said little. President Zelensky, who was not there, called the talks "constructive and quite substantive". US Secretary of State Marco Rubio said they were "very productive".

Few expect a breakthrough soon. US officials say both sides are showing "a little bit of flexibility" on who will control eastern Ukraine after the war. Prisoner swaps were also discussed. European leaders have mostly been left out, but the US now says they are part of the answer.

WHY IT MATTERS: *These are the main talks on ending a war that began in 2022.*

BBC News · Nov 10

STRIKE

Spain holds its first general strike in 14 years over high rents

Millions of workers in Spain stopped work on Wednesday to demand action on housing. The two main unions, UGT and CCOO, called the strike. Buses and trains ran late or stopped in Barcelona and Madrid, and many schools closed or ran at low levels.

The unions said 90% of workers in the car and port sectors joined, and 75% in energy. Rents in Spain have almost doubled in ten years. The strike comes less than three weeks before a national election.

WHY IT MATTERS: *It shows how much housing costs now shape politics in Europe.*

NPR (Associated Press) · Nov 11

HEALTH

Congo's Ebola outbreak passes 11,000 cases

Congo's government says 11,042 people have caught Ebola since May,

and 5,387 have died. It is the country's deadliest outbreak. The virus, called Bundibugyo, has no known vaccine or treatment. It has reached 8 of Congo's 26 provinces.

Fighting in Ituri and North Kivu provinces makes it hard to reach the sick. Some health workers have been attacked by people who distrust the virus warnings. The disease has also spread to Uganda, Kenya and Nigeria.

WHY IT MATTERS: *The worst Ebola outbreak so far, in West Africa from 2014 to 2016, had about 28,000 cases.*

NPR (Associated Press) · Nov 11

KENYA

Kenyans in Nairobi's river slums fear floods from El Niño

El Niño is a warming of the Pacific Ocean that changes weather around the world. Forecasters say this one could be among the strongest on record. In March floods killed more than 60 people in Kenya.

Crews are clearing rubbish from Nairobi's rivers and drains. The city's drains were built 30 years ago for 2.5 million people, and about 7 million are there by day. A full upgrade would cost up to $192 million, which the city lacks. Officials say up to 100,000 people could be at risk in the worst case.

WHY IT MATTERS: *Poor neighbourhoods by rivers take the worst of the flooding.*

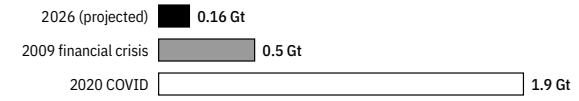
BBC News · Nov 11

REGION FIVE

Science

Emissions dip for a reason that may not last, and fungi that hold up grasslands.

HOW BIG IS THIS YEAR'S DIP IN CO2?



Science News, Nov 10, on a Nature Climate Change study

CLIMATE

Carbon dioxide from fossil fuels is set to fall slightly in 2026

Scientists project emissions will fall 0.4% this year. From 2019 to 2025 they grew 0.8% a year on average. The study was published Nov 9 in Nature Climate Change.

The likely cause is high oil prices from the Iran war, not a faster switch to clean energy. Electric cars, wind and solar are helping, though. If they keep growing, they could cancel

about 63% of the growth in emissions from 2025 to 2028.

WHY IT MATTERS: *A dip caused by a price shock may not last.*

Science News · Nov 10

ECOLOGY

Hidden fungi may help save Mongolia's dry grasslands

More than 70% of Mongolia's land is damaged by heat, drought and overgrazing. The country is warming three times faster than the world average. Researchers want to map the underground fungi that trade nutrients with plant roots and help plants live in poor soil.

The idea is early. The team is still looking for the places where these fungi are most common, so those spots can be protected first.

WHY IT MATTERS: *Protecting the soil's hidden life might be cheaper than replanting land that has died.*

Nature · Nov 10

BUBBLE CHECK

The three newsletters that arrived since last night are about the midterms, India's economy and how to take criticism. None is about AI. None of today's news stories is either. For something different, read the Spain story: the price of a flat can move a whole country to strike.

READ SLOWLY

The Essays

Printed whole, in the writers' own words

ASTRAL CODEX TEN · NOV 11 · OPINION

Escaping The "Outside Criticism" Trap

by Scott Alexander · Astral Codex Ten

A psychiatrist-blogger explains a trick used to call anyone closed-minded, and how to answer it.

Here's a fun argument you can use to discredit anyone at any time. Start by calling them an arrogant fanatic with a totalizing ideology, someone who thinks they personally have solved all problems and can't possibly learn from dissenting viewpoints. If they have friends who agree with them, they're a cult who collectively work to shelter themselves from any dissenting viewpoint.

Then your victim responds by presenting one hundred notarized affadavits describing times when they have engaged with their critics, spent years studying books by their critics to see what they could learn from them, and even changed their beliefs in response to criticism.

Then you spring the trap: "Ah, but you cherry-picked critics who operate within your worldview, a sort of loyal opposition. You continue to dismiss and exclude genuine critics who challenge your smug assumptions."

This always works. If your victim protests that they do like and appreciate many critics with worldviews different to theirs, you object that no, those guys are also within the system. You meant a different guy, even further away from their own view. Oh, they read that one too? You meant an even more different guy. Eventually, no matter how cosmopolitan they are, you will reach someone they actually don't respect - let's say a Stalinist pedophile who believes the Papacy is a Jewish plot - and then you go ha! I told you that you rejected everyone outside your little bubble!

Everyone has people close to their ideology who they respect, and people far away from it who they don't respect. Professionally, I work from a tradition of medicalized psychiatry influenced by figures like Kraepelin, Kline, and (sigh) Stahl. When I operate within that tradition - let's say by prescribing escitalopram to a depressed patient - here are five circles of critics I might consider listening to:

The Colleague: "Instead of prescribing escitalopram, you should prescribe bupropion, which studies show is equally efficacious with fewer side effects."

The Biopsychosocialist: "Instead of just using medication, you should treat depression with a balanced regimen of biopsychosocial interventions, including improved diet, better sleep schedule, and helping the patient overcome obstacles in their life."

The Freudian: "Have you even considered that the patient might be repressing his latent desire to kill his father and marry his mother? Maybe you should analyze him weekly for ten years."

The Zealot: "Mental illness is caused by separation from Christ/Mohammed/Buddha/Confucius. Your patient should seek forgiveness and let God/Allah/enlightenment/filial piety into their heart/soul/chakras/family."

The Scientologist: "Psychiatrists are hostile aliens from the planet Farsec. Mental illnesses are caused by body thetans under the influence of the evil overlord Xenu, which can only be removed through dianetics classes, especially the expensive ones."

In practice, I devote a lot of time and energy to the Colleague and the Biopsychosocialist's criticisms, a small amount to the Freudian's, and very little to the Zealot's and Scientologist's.

If you portrayed me as a blinkered fanatic who refused to consider other views, I might object by talking about all the time I spent engaging with my colleagues and the biopsychosocialists. Then you could say that those are just loyal opposition within my own modern-20th-century-medicalized-psychiatry paradigm, but, if I'm so broad-minded, why don't I engage more with Freudians, zealots, and Scientologists? I can only answer that I think those people are wrong. Then you could say "See, , Scott admits that he thinks that his own ideology is the sole source of truth and that outside critics have nothing useful to contribute".

This is what Eric Raymond calls a Kafkatrap, where any denial is further evidence of guilt. The victim insists "I do so listen to criticism! See, I draw my boundary of who to listen to way out here", and points to some place very far away from him. The attacker answers "See! Everyone inside that boundary is within your worldview by definition, so you're admitting that you refuse to listen to criticism from outside your worldview!"

Raymond coined "kafkatrap" during the first round of Social Justice Wars, and this argument figured heavily there. You're in medicine? You read dozens of books by highly-educated critics of medicine all across the political spectrum? Ah, but have you engaged with this anthroplogist who says medicine is fundamentally racist because surgeons don't start every operation with a land acknowledgment? Oh, you don't take her seriously? Seems

like you're one of those insular elitists who think true outsiders, especially black outsiders, have nothing to contribute. If you get my drift, mein Fuhrer.

This is also part of what I was getting at in my plea for effective altruists to stop fawning over their own worst critics. I think rationalists and effective altruists have done a great job engaging with outsiders who have helpful perspectives on things they care about. But the demand - then and now - has always been that we go further. Why aren't we engaging with the people who think charity is evil? Why aren't we engaging with the people who think "reason" is a racist scam to claim that white values are objectively correct? Why aren't we bowing and scraping before the people who spent the past five years saying we were total idiots to worry about AI because it would never be able to count the number of r's in "strawberry"? Why don't we worship them as gods? Why don't we reach out to the people who say we are ugly pervert techbro nerds who should kill themselves, and offer them all-expenses-paid trips to give the keynote speeches at our conferences?

When I try to dialogue with the people who say these things (they're critics! I have to dialogue with them), I try to ask them what level of "engagement" would make us no longer insular and arrogant and exclusionary - what would allow multiple people in the same social network to share an unpopular belief without getting called "cultish"? I never get a straight answer, and I suspect this is because the real answer would be "You have to abandon all of your unpopular beliefs and switch to believing only popular things instead, and also make me and my friends the high-status leaders of your group". I decline to do this.

Granting that these people suck, and stepping back from the kafkatrap for a moment, how should we address this epistemic problem? Certainly as a first approximation you shouldn't spend too much time on people you think are dumb and crazy. But anyone with an opinion sufficiently different from yours probably seems dumb and crazy. If we never listen to anyone whose opinion seems too different from what we currently think is correct, how can we learn?

There's one very tempting way out: look for "objective evidence" of which critics are worth engaging with. Someone who is very popular and powerful and well-regarded needs to be taken seriously; some marginalized group who can't get their ideas published in high-impact-factor journals doesn't. Biopsychosocialists have more social cachet and published journal articles than Scientologists; this accurately predicts that they should be taken more seriously.

But there are two flaws in this plan. First, your perception of "powerful", "popular", and "well-regarded" is probably influenced by your current be-

liefs. I mentioned "high impact-factor journals" above - why do I think those are better "objective" evidence than the number of mentions your ideology gets in Battlefield Earth, one of the best-selling science fiction novels of all time? Probably because I'm a doctor and not a Scientologist.

Second, and more important, it also ensures you always move in a gradient towards Power. Every step you take will be away from your group's own values, in favor of consensus society. "Abandon everything that makes you unique in favor of being exactly like everyone else" is a Heuristic That Almost Always Works. But it's also the sort of thing demons whisper to you on desert roads at night. It would be wise to have a plan that tells you something different from the demons.

I rely a little on objective evidence, but I also think it's healthy to do something like exponential decay: halve the amount of energy you devote to each step away from your current beliefs. If you think medicalized psychiatry is the way to go, you should spend most of your time trying to be the best medicalized psychiatrist you can. Then spend half that amount of time investigating things in the near periphery of medicalized psychiatry, like cognitive therapy and biopsychosocial solutions. Then spend half that amount of time on things in the very far periphery, like Freudianism. Then spend half that amount of time mopping up the things that seem obviously crazy.

"Half" is just an example (I personally do more like a fifth). It doesn't even have to be the same constant at each step. I think the key idea is that instead of beating yourself up because you don't read people you hate at exactly the same level as people you respect, you can just read them a little but not too much. This has the advantage of helping you check whether they're actually as silly as you thought, or whether you were actually missing something important.

(The people trying to kafkatrap you will insist that you didn't truly understand the few books you read, and that you need to spend years truly soaking yourself in the alternate paradigm before it will can start to seep into someone as beknighted and biased as yourself. This . . . might be true, but there's no way to spend years investigating every idea, and you shouldn't award years of your life to whoever can bully you the most, so I still think concentrating on correct-seeming ideas and devoting progressively less energy to progressively more outlandish-seeming ideas - and hoping that this small amount of energy and study will be enough to help you catch the good ones that you should spend more time on - is the way to go.

■

astralcodexten.com/p/escaping-the-outside-criticism-trap · Scott Alexander is a psychiatrist who writes about science and rationality; a list of book reviews at the end is left out

IN PLAIN WORDS, AND THE OTHER SIDE

The short version: you can attack anyone as closed-minded. If they show they read critics, you say those critics were too close to them. Eric Raymond called this a "Kafkatrap", because every answer counts as proof you are guilty. Alexander says you should not get lost in it. His rule is to spend most of your attention on views near your own, and a bit less on each step further away, maybe half as much each step. That way you do read outsiders without feeling you must give every view equal time.

The other side. This is one writer's view, and he is a target of this kind of criticism, so he has a stake. Critics say "most time on views near mine" can be a polite way to ignore people who disagree the most, and that those are often the ones who spot your blind spots. He answers that you can't read everything, and that famous or powerful critics aren't always the right ones. The essay gives no test for who is worth reading.

BENTHAM'S NEWSLETTER · NOV 10 · OPINION

Longtermism For Animals

by Bentham's Bulldog · Bentham's Newsletter

A philosopher-blogger argues that the animal movement should focus on how animals fare in the far future.

If you're serious about helping animals, you should be almost entirely concerned with making sure the far future goes well for animals.

The core argument for Longtermism is that the future is huge! There could be countless future generations that our actions might affect. Thus, in expectation, for every one present individual affected by your actions, there are countless affected future individuals. If you want to maximize your positive impact, you should care a lot about the impact you have on the future—because that is most of the impact that you have.

This argument applies to animals just as much as it does to humans. In expectation, vastly more animals will live in the far future than live in the present day. If we're concerned about animals, we should be very concerned with making sure the future is good for animals. That's even more important than stopping present animal suffering. An animal being tortured in a factory farm a thousand years from now is just as tragic as one being tortured in a factory farm today.

Because the future is so huge, taking it seriously has pretty nutty implications. For example, one extremely conservative estimate found:

> *Reducing the probability of spreading [wild animal suffering] to one other planet by 0.1% has an expected value equivalent to stopping all suffering caused by exploiting animals for food on Earth for over a thousand years.*

My guess is that this estimate is too low by many orders of magnitude. But the basic logic is straightforward: if we bring animal suffering to a new planet, there could be suffering until some cosmic catastrophe ends life on that planet. This could mean billions of years of suffering. Actions that nudge the probability of spreading suffering to other planets almost totally dominate actions on factory farming (though ending factory farming is one good way of stopping future animal suffering).

And note: animal Longtermism is primarily about stopping extreme suffering in the future. It is a lot more obvious that preventing future extreme suffering is important than it is that creating lots of happy people is important. So the premises on which the case for animal Longtermism rests are pretty conservative.

The future is big enough that if the animal movement wants to make things go well for animals, they can for all practical purposes ignore what will happen near-term, except insofar as it affects the probability of things going well long term. The short-term impact of the animal movement is likely to be basically non-existent compared to their impact in the very long term.

So what is it that those in the animal movement should do?

First, it's really important that we don't spread factory farming or wild animal suffering to space. If you think animals live bad lives, making sure we don't fill countless planets with the screams of innocent terrified animals is hugely important. Preventing intergalactic factory farms is a no-brainer.

Second, we need to make AI concerned about animal welfare. We are currently building gods of metal; it is reasonably likely that they'll be indifferent to almost all the well-being and suffering in the world. This is partly for the sake of animals. AI is likely to soon become extremely powerful. Making sure that the digital beings who will dominate the future aren't indifferent to animal welfare is of vital importance.

My friend Adria Moret has a piece written alongside Peter Singer and two other people that makes this case in more detail and talks about how it could work. In my view, getting AI models to have concern for animals is the #1 near-term priority for animal welfare. If I could trade a doubling of factory farming for getting all AIs firmly and reliably committed to its abolition, I would take that trade. Kudos to both OpenAI and Anthropic for including concern for animals in their model specs.

There's also a more self-interested reason to make AI concerned about animal welfare: if AIs internalize norms that neglect the interests of alternative species that are much dumber than them, then they would neglect our interests once they become superintelligent. It is thus important that AIs

take every sentient being seriously. If you want to get more involved, I recommend attending Sentient Futures conferences and applying to work for organizations like Sentient Futures or the NYU Center for Mind Ethics and Policy. The welfare alignment project, in particular, is thinking through AI alignment for animals in a very serious and responsible way.

Third, trying to increase compassion for animals—especially neglected animals like shrimp and wild animals—is likely to make the future go well for animals. A world that cares more about the welfare of animals is less likely to harm them gratuitously without a good reason.

This also supports efforts to work on particular animal welfare causes that might stick. So, for example, I think Screwworm Free Future is a great charity to support. Screwworms are horrible—they lay eggs inside live animals, so that little maggots devour their victims from the inside. Screw screwworms! Here's , who pioneered the use of CRISPR in gene drives, summarizing the situation:

If we're serious about wanting to benefit animals, we should mostly be concerned with how our actions affect them over long time scales. If your animal welfare intervention has no chance of making things better more than a hundred years out, you should probably stop doing it. Fortunately, a lot of standard animal welfare interventions push the world in a more pro animal direction, making it likelier we'll get a future that's good for animals.

The scale of future animal suffering is truly incomprehensible. This is one of the big challenges that we'll need to solve to get a good future, and it's the most important challenge in the animal welfare space. For this reason, if you're an animal advocate reading this, I cannot implore you enough to think hard about the far future.

■

benthams.substack.com/p/longtermism-for-animals · two upsetting passages about parasites and a link list are left out

IN PLAIN WORDS, AND THE OTHER SIDE

The short version: "Longtermism" is the idea that the future could hold so many lives that helping them matters most. The writer says this also applies to animals. If humans spread factory farming, or wild animal suffering, to other planets, the harm could be enormous. He says animal groups should work to stop that, to make future AI care about animals, and to make people kinder to animals.

The other side. This is one writer's view, and the numbers in it are guesses about a very far future, as he admits. Critics say effort spent on distant fu-

tures can take attention from animals suffering now, in farms today, where we know what helps. Others doubt that anyone can predict what will matter for animals in thousands of years. The writer is a committed longtermist, so this is the case from inside that view.

YOUR INBOX

The Newsstand

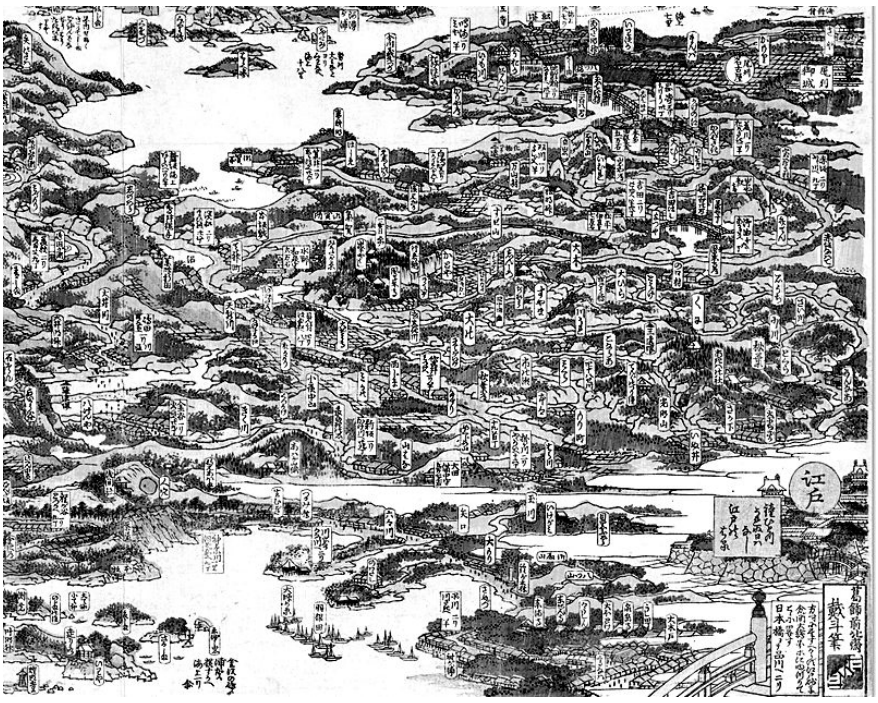
The last 24 hours, summed up

NOAHPINION · Noah Smith

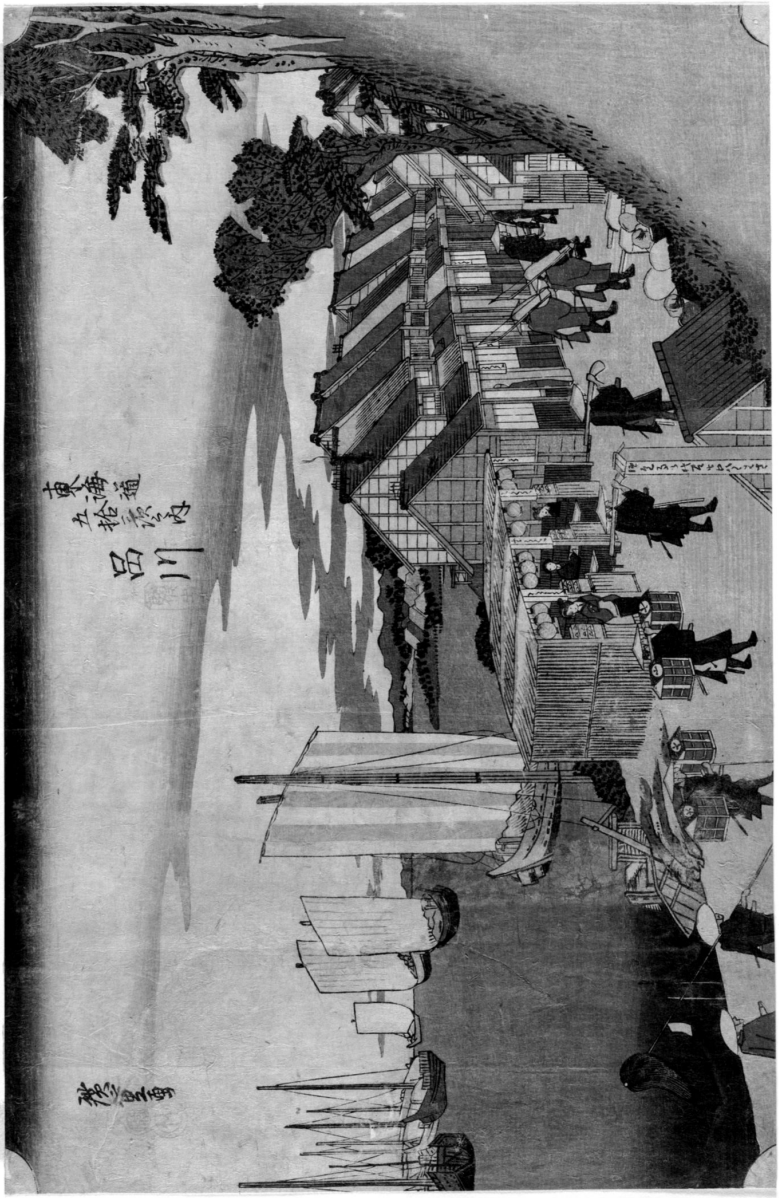
SUMMARY

India is quietly doing great

Opinion, from Wednesday. The writer says India is doing very well. Extreme poverty has mostly ended, the economy grows about 6.5% a year, and the government is building roads and trains fast. He says it may become the brightest big economy. He names open problems: weak private investment, poor rural schools and few women in work.



Katsushika Hokusai, Famous Places on the Tōkaidō in One View (c. 1818). UBC Library, no known copyright restrictions.



THE ROAD · STOP 2 OF 55

Shinagawa: Sunrise in the Fifth Month · Hiroshige, c. 1833–34

Shinagawa was the first town travellers reached after leaving Edo. It sat on the shore of Edo Bay, and its inns, teahouses and boats are all in this view. Hiroshige made 55 prints of the road, and this paper shows one each issue. Turn the page sideways to see it the right way up.

Public domain (CC0). Cleveland Museum of Art, via Wikimedia Commons.